



Test your Architectural Resilience with Chaos Engineering

Marcel Neidinger

Solutions Architect
Amazon Web Services



"Everything fails, all the time"

Werner Vogels

CTO Amazon Web Services

ed the Turing tes

A few definitions

Resiliency

The ability for a system to recover from a failure induced by load, attacks, and failures.

Availability

Percentage of time a workload is available for use.

Fault Tolerance


Handle interruptions without degrading the system performance

An Introduction to chaos engineering




Different ways of testing


Unit Testing



Scenario Testing*

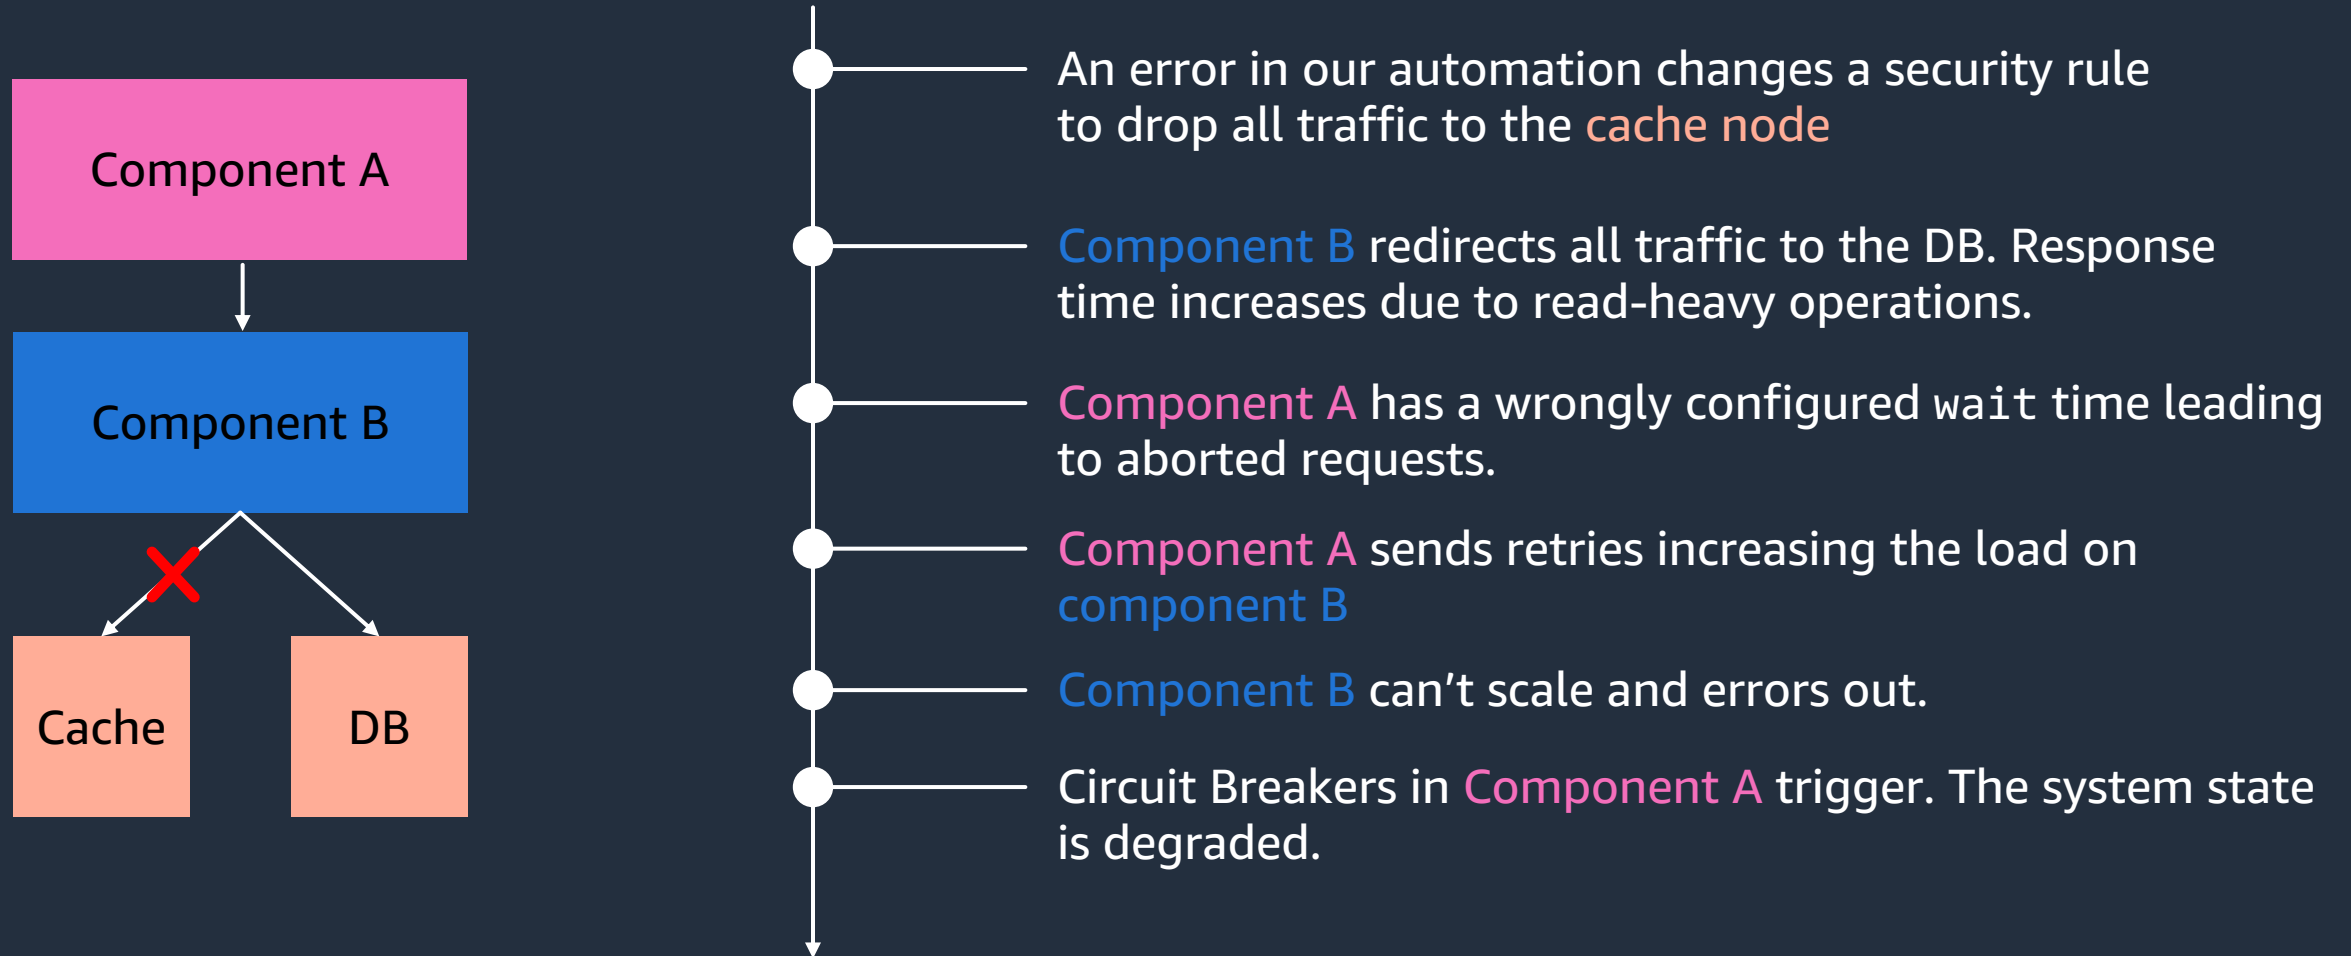


(Chaos) Experiments



Failures are usually chain reactions of small incidents

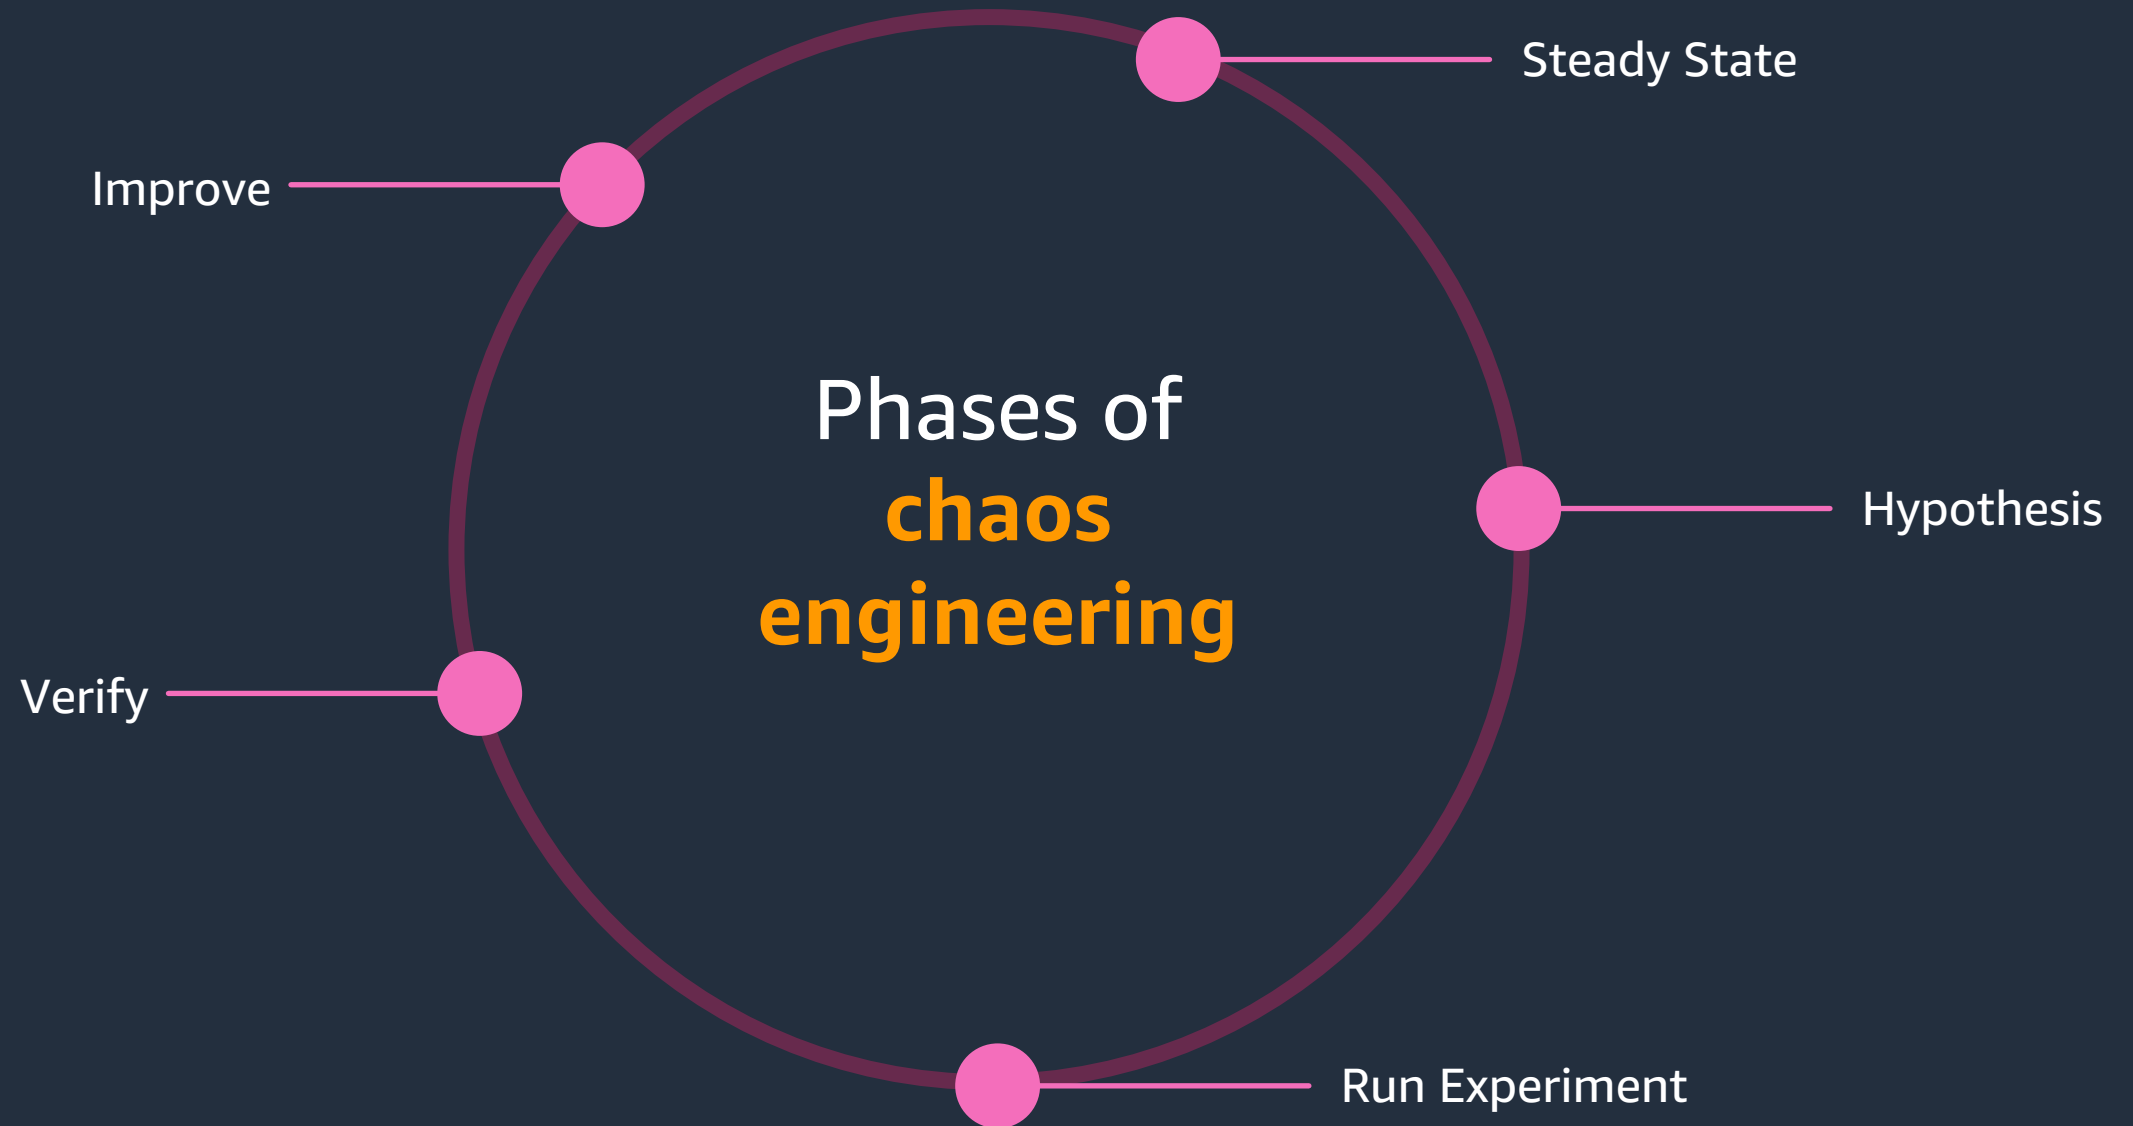
A TIMELINE

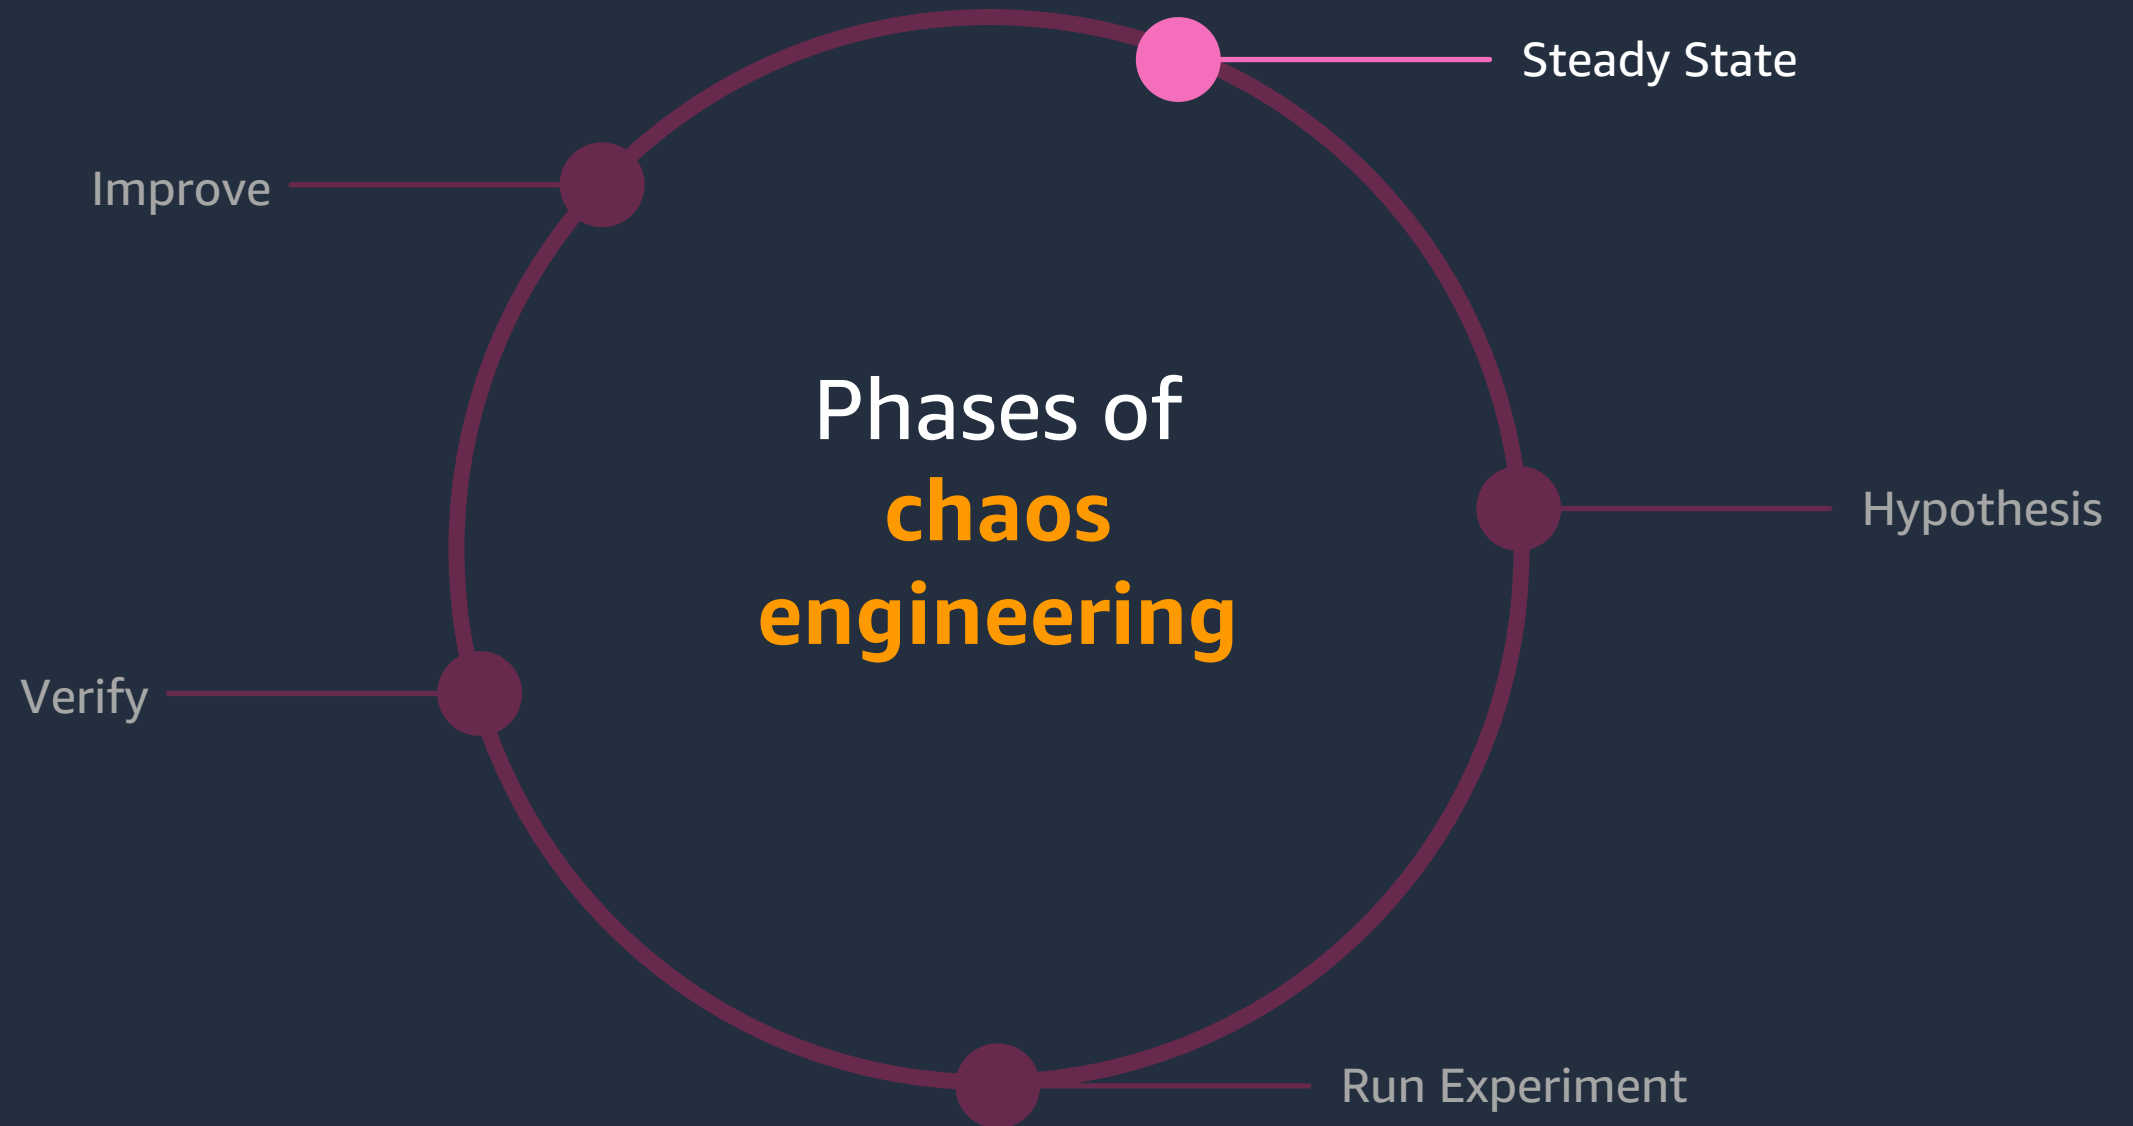


Chaos Engineering is **NOT** about **breaking things** randomly **without a purpose**.

Chaos Engineering **is** about breaking things in a **controlled environment** and through **well-planned experiments**.


While you **CAN** do this in production you **DON'T** have to.





Steady State


THE “NORMAL” BEHAVIOR OF YOUR SYSTEM



Business + Ops

Observability as a Pre-Requirement

HOW DO YOU MEASURE “NORMAL”?

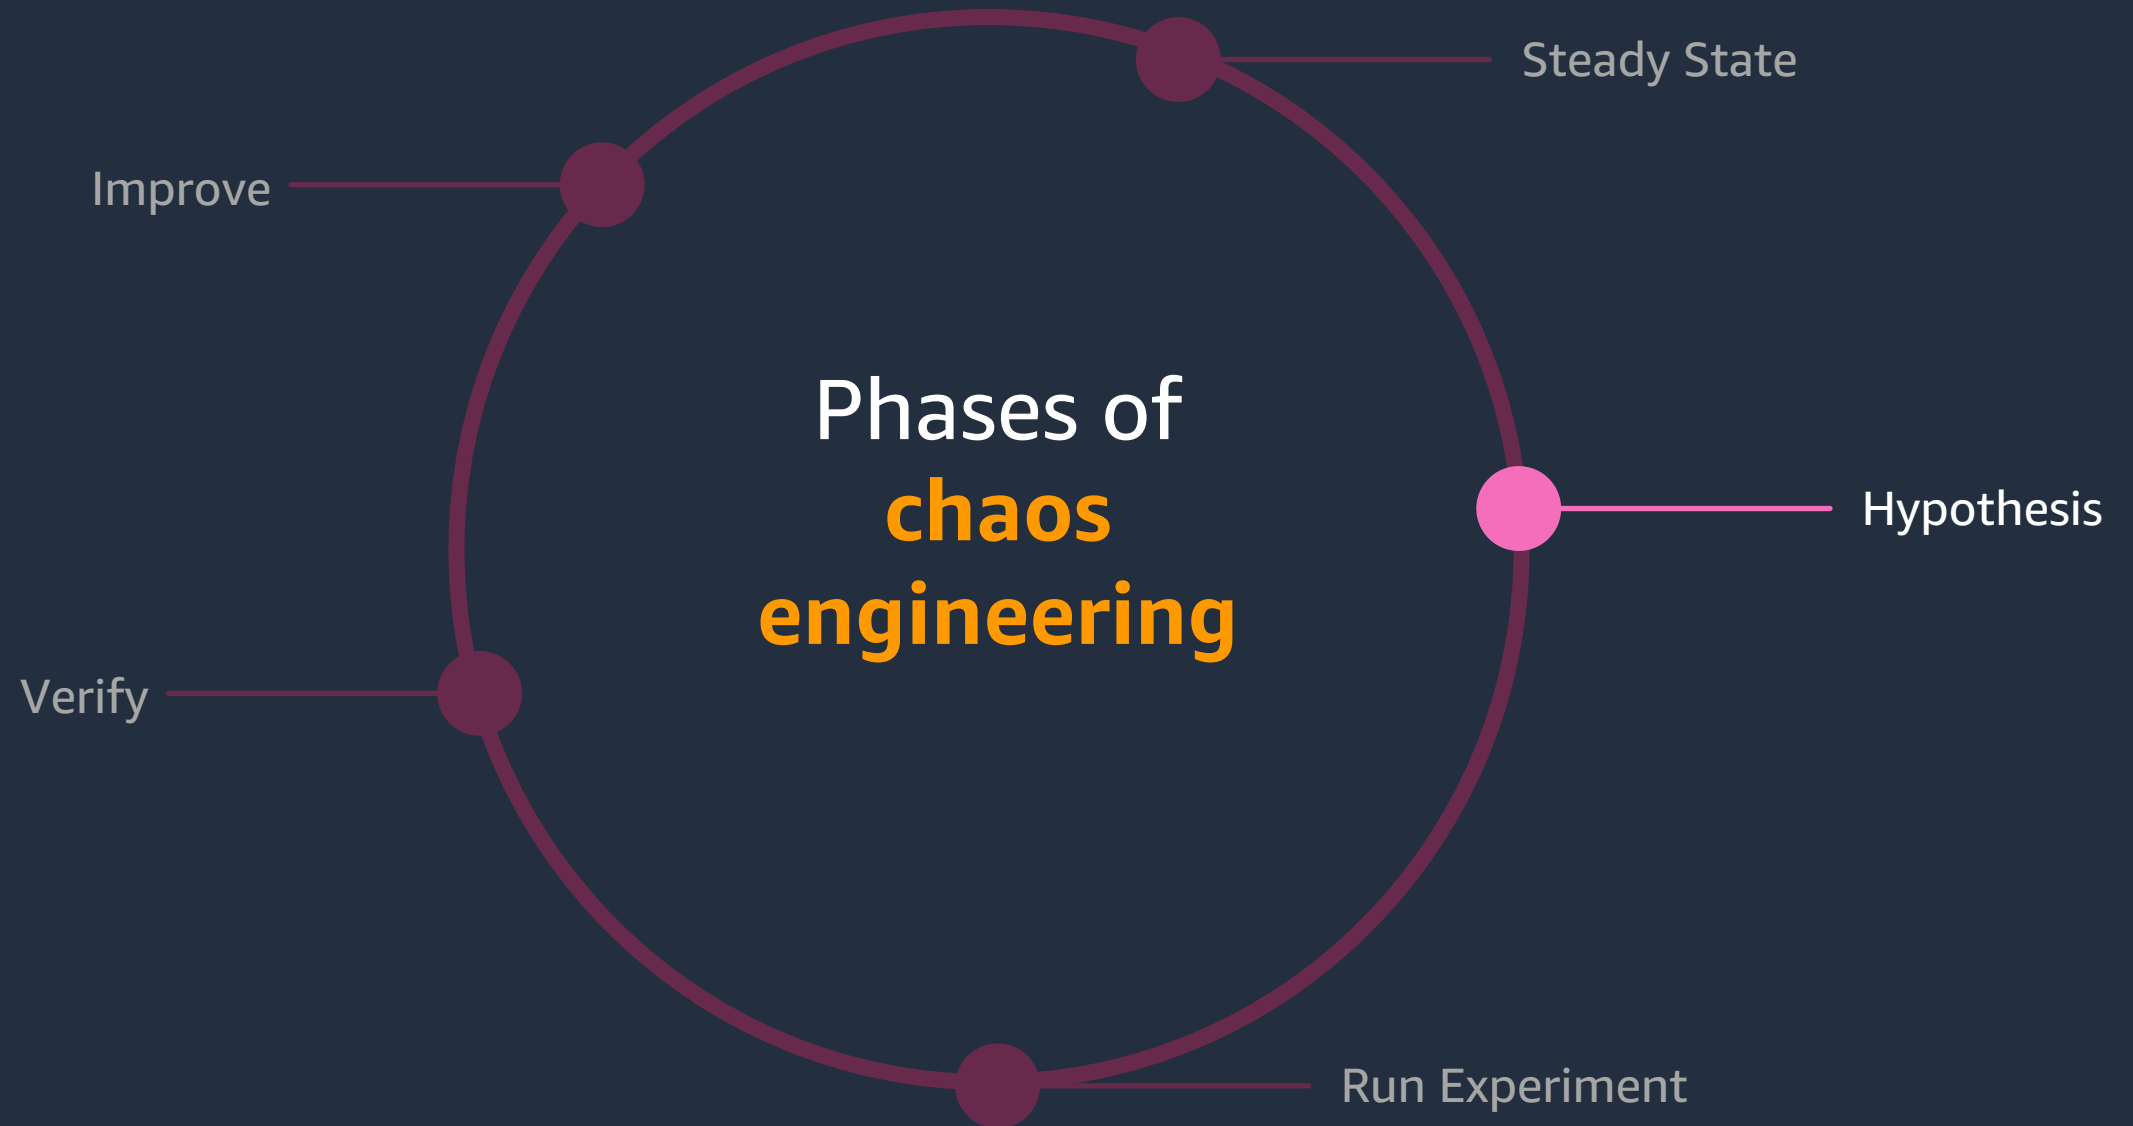


Do you understand the **inner workings** of your application?

Do you understand **any system state** your application may have fallen into?

Can you understand the above, just by **observing** your tools?

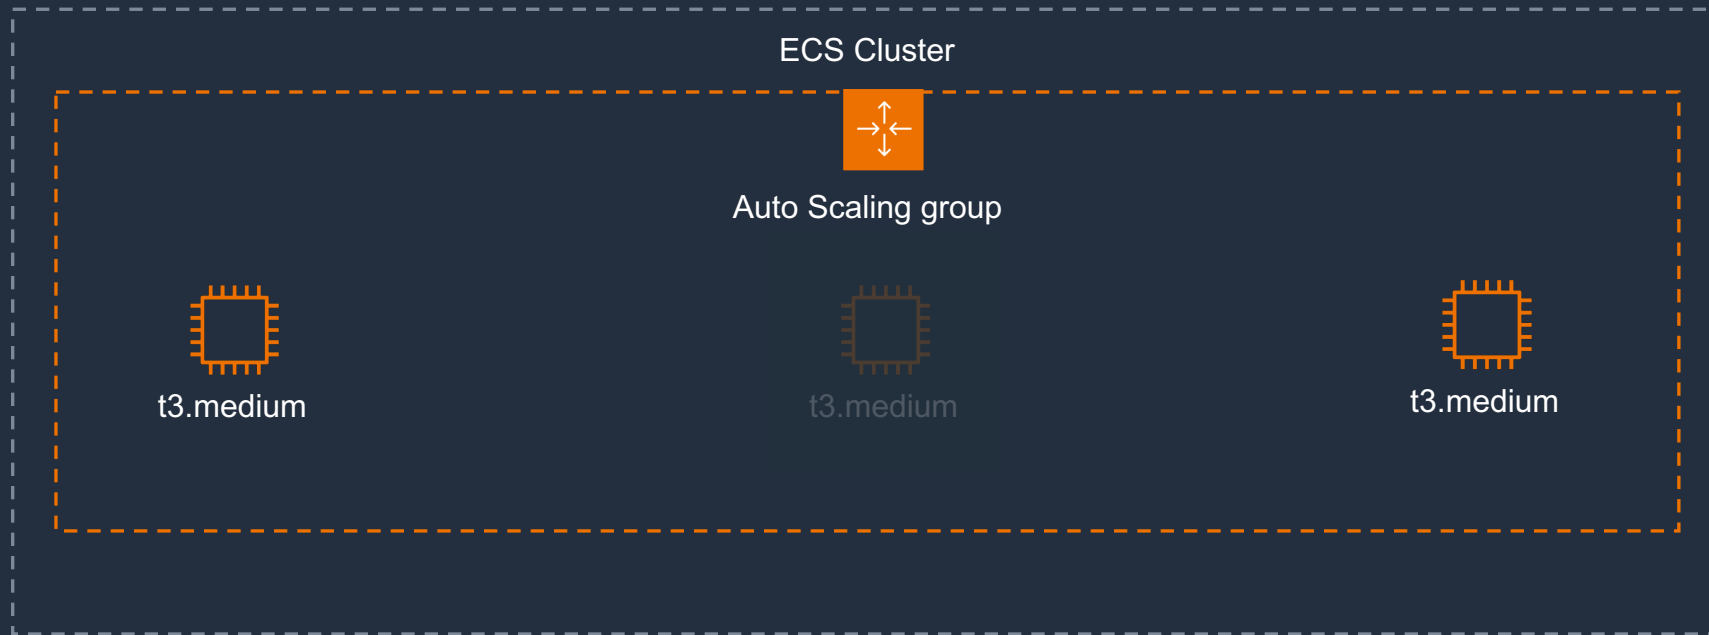
Can you understand the state, no matter how **unusual** the situation is?



Hypothesize that this **steady state** will **continue** in both
the **control group** and the **experimental group**.
unaffected parts of our system *affected parts of our system*

Hypothesis

What if a host of our ECS cluster shuts down?

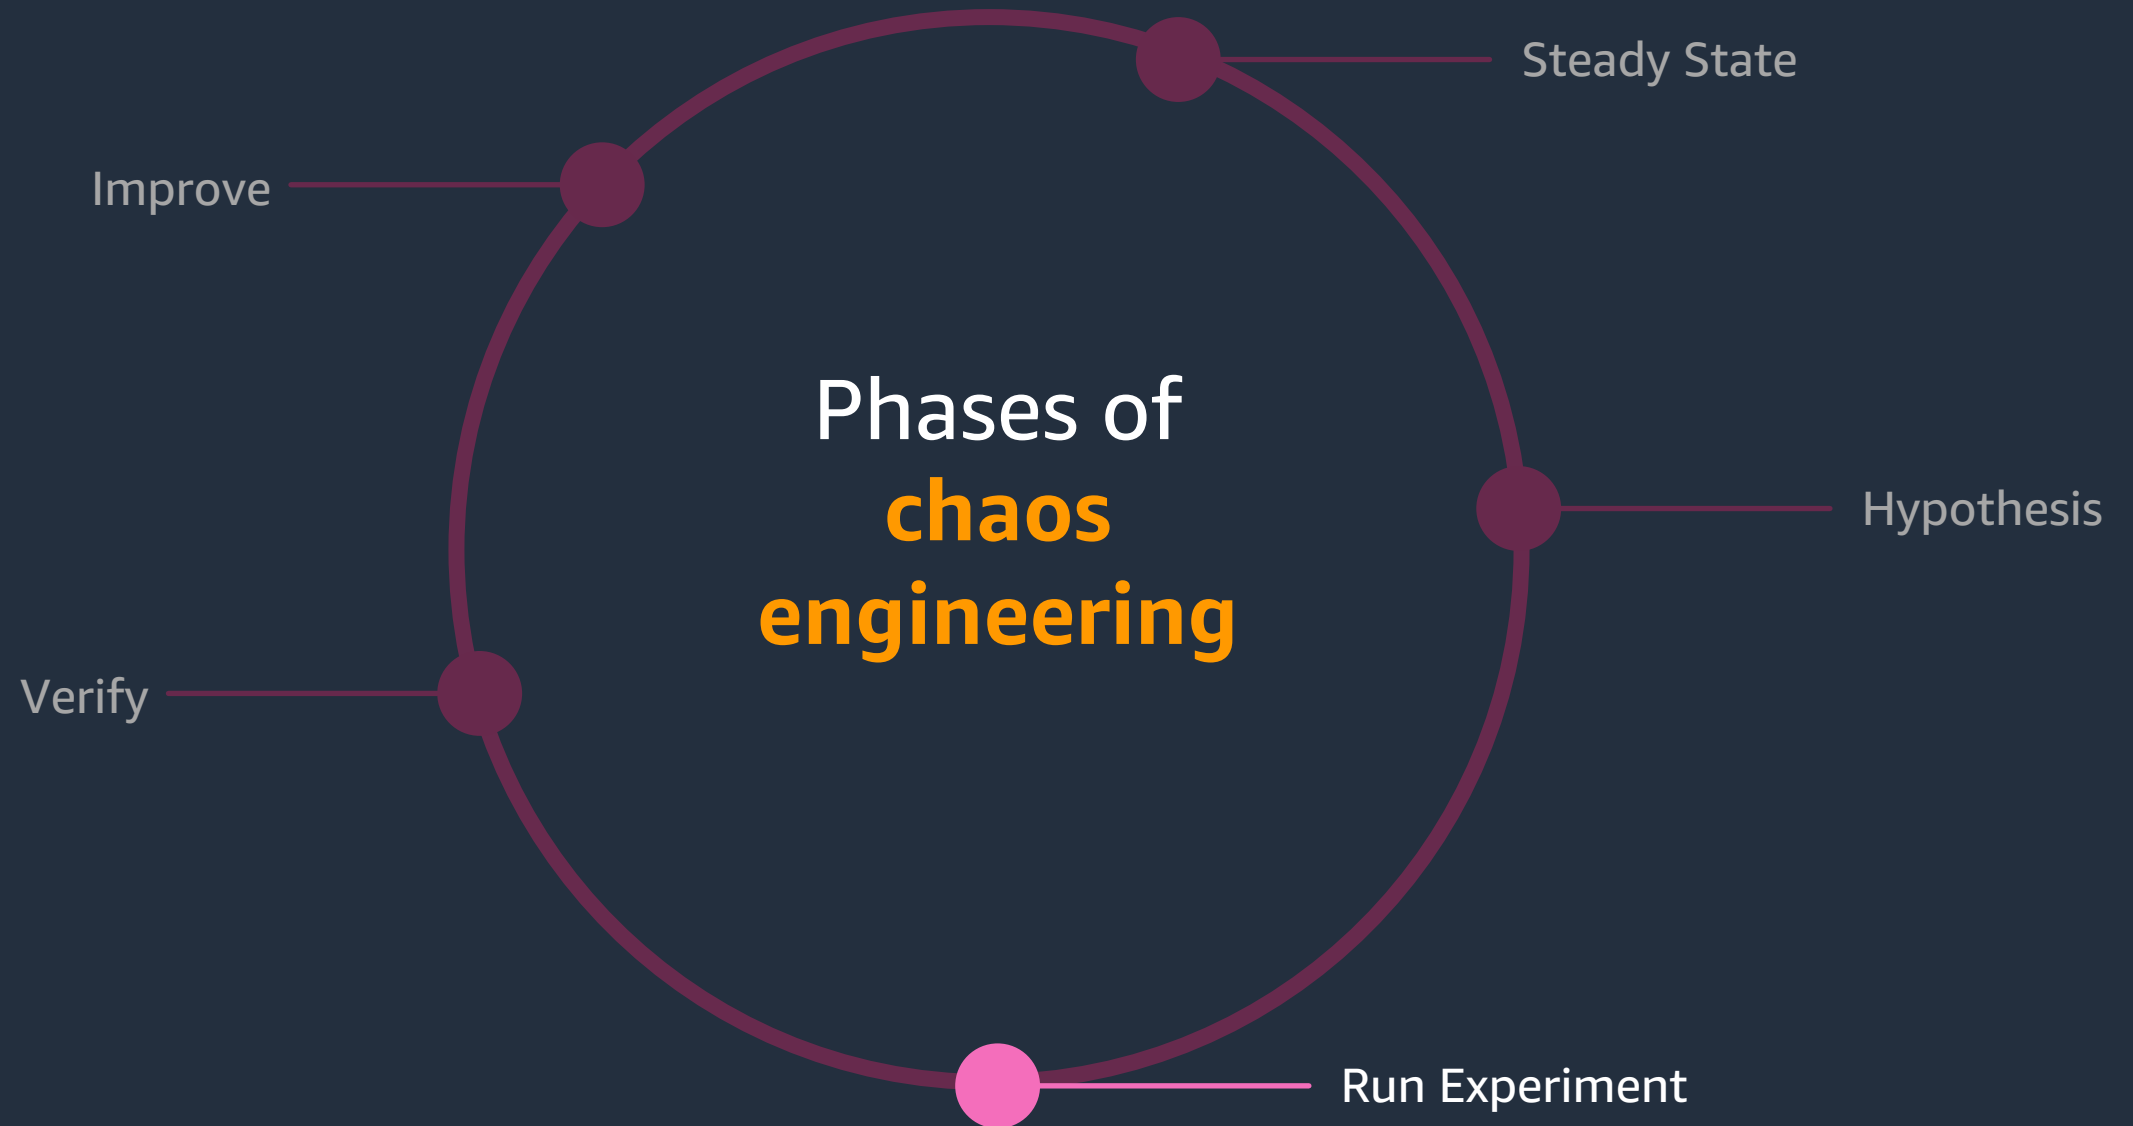


Hypothesis Considerations

What if a host of our ECS cluster shuts down?

Have we architected against this?


Do we want to guard against this?




Designing real-world experiments



Code & configuration
e.g., bad deployment,
cred expiration,
host shutdown



Infrastructure
e.g., datacenter failure,
hardware failure



Data and state
e.g., data corruption,
overload




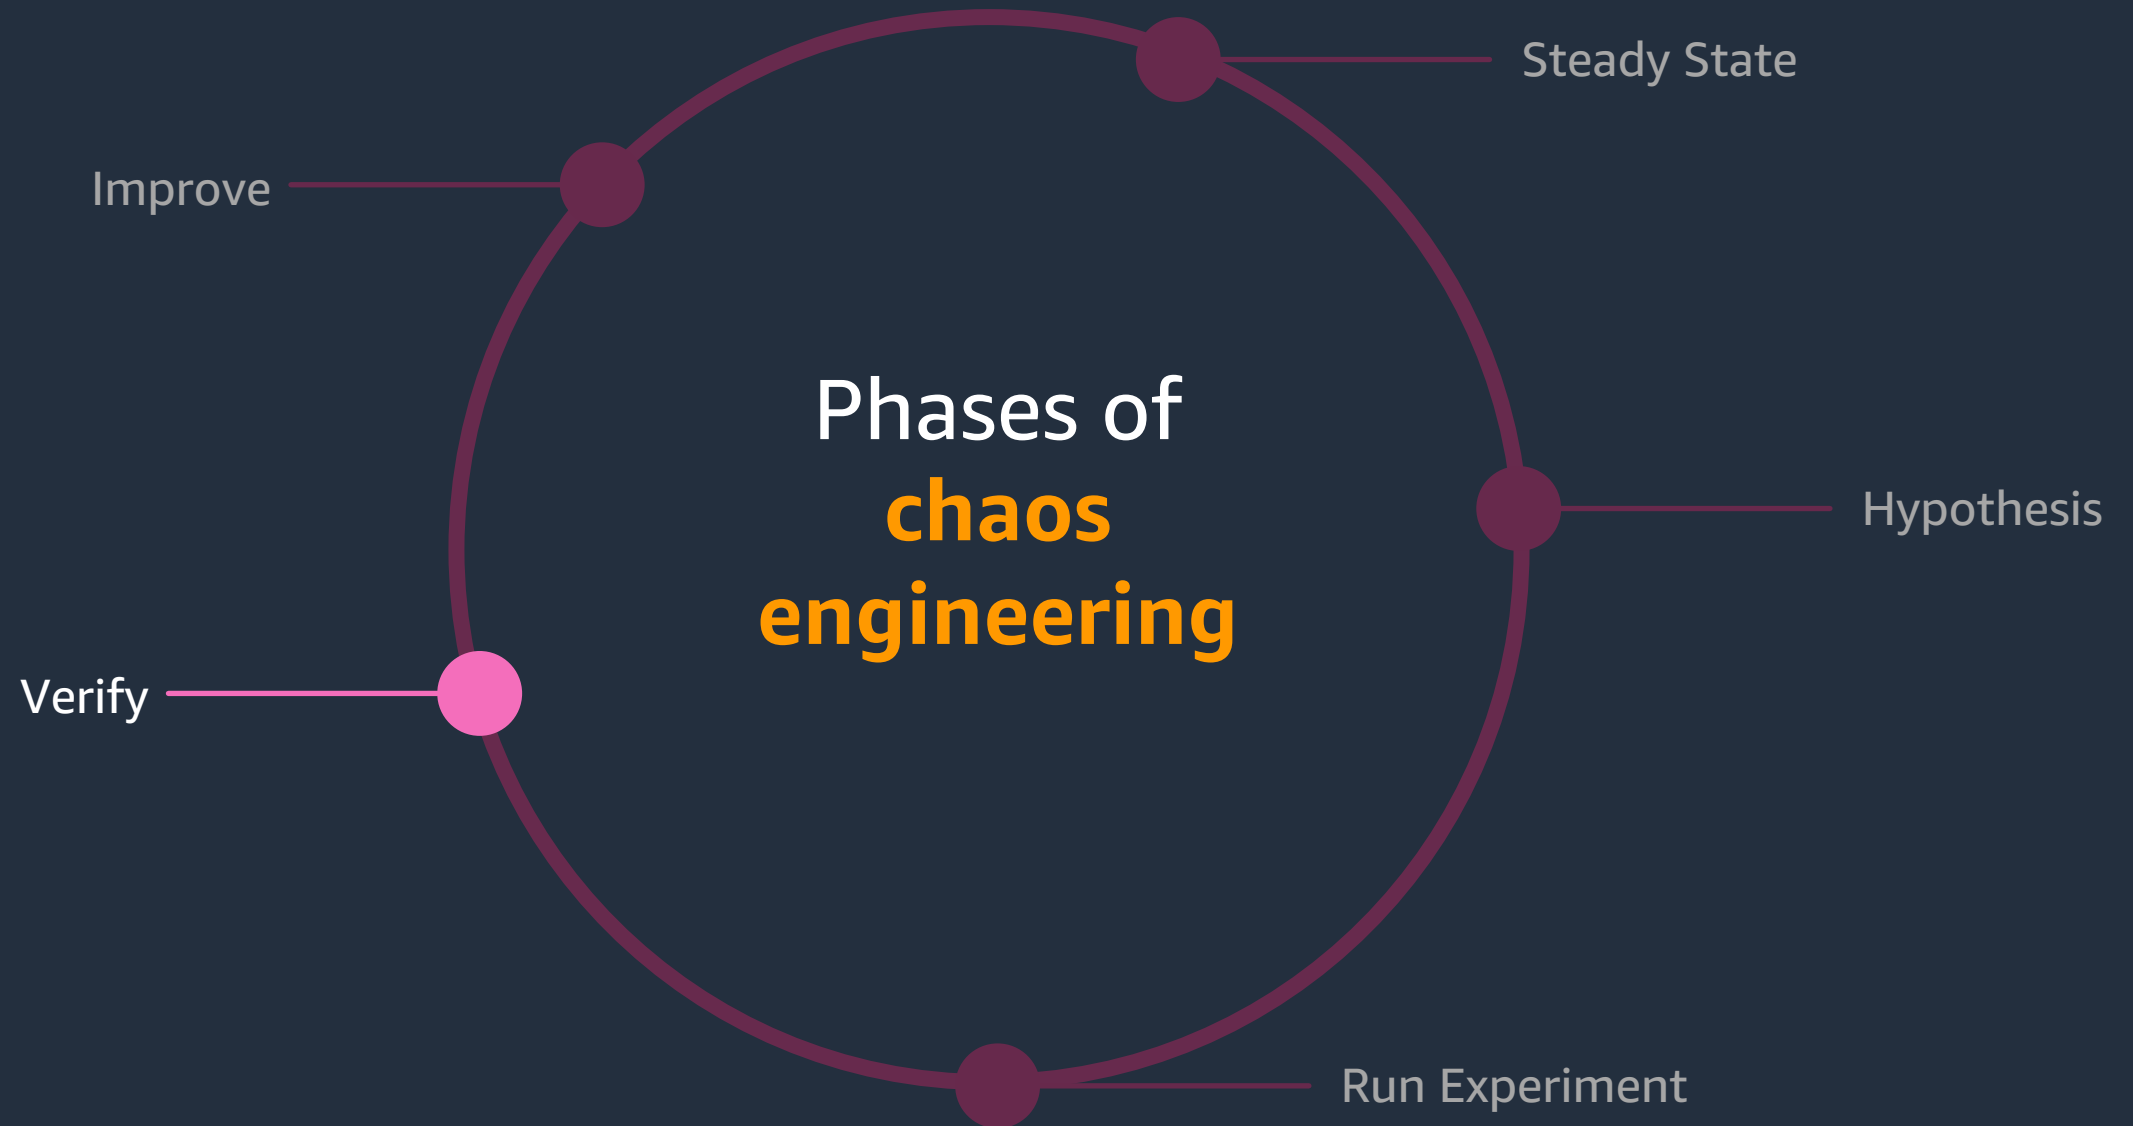
Dependencies
e.g., third-party
integrations, AWS services



Highly unlikely, but technically feasible
e.g., physical loss of an entire
Region, the internet is down

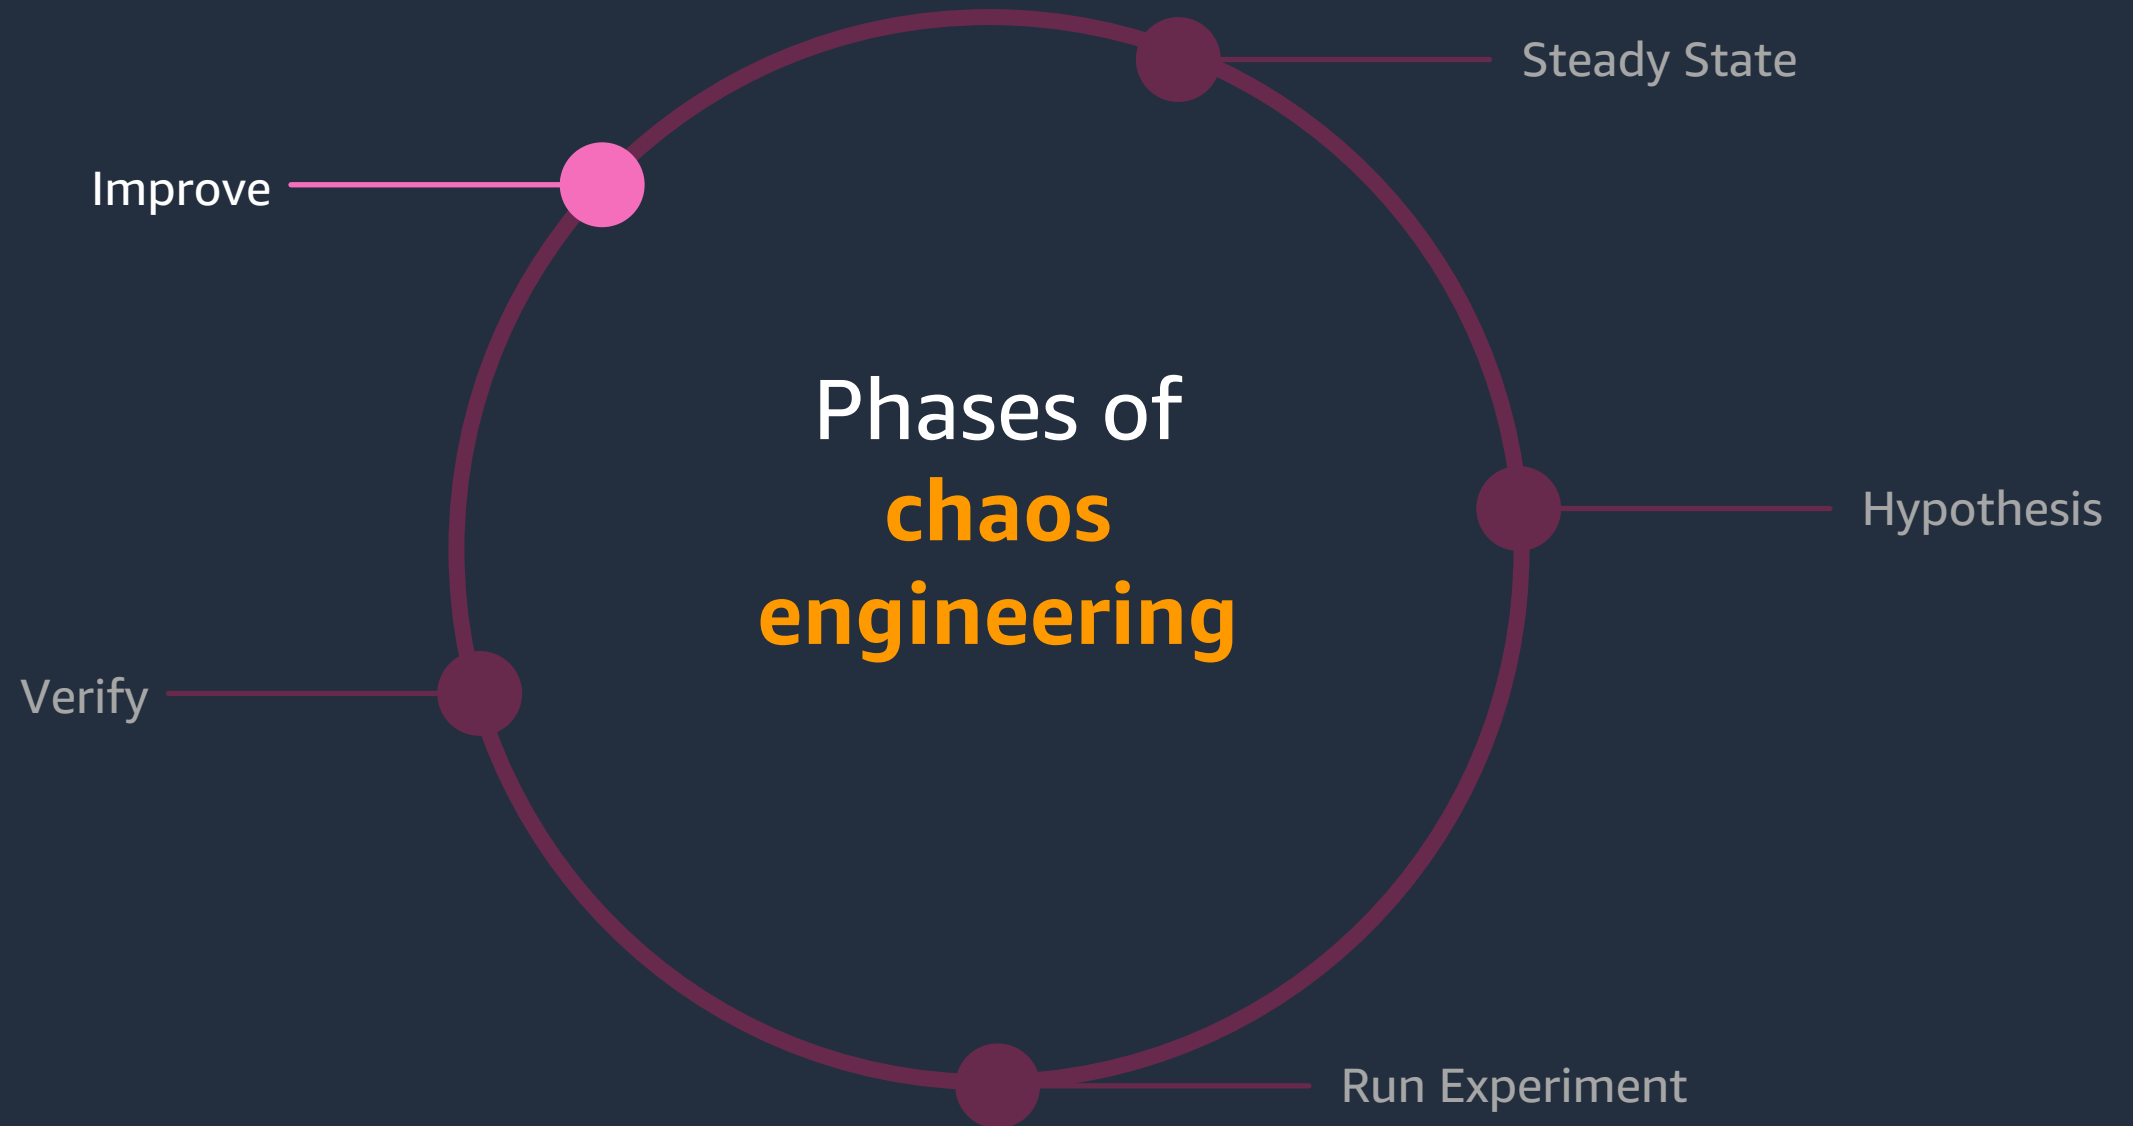
Start small





Quantify your experiment

- Time to detect?
- Time to notification?
- Time to escalation?
- Time for graceful degradation to start?
- Time for self-healing to happen?
- Time to (full|partial) recovery?
- Time to all-stable?



Fix it!

A simple example of chaos engineering



AWS Cloud



eu-west-1



Virtual private cloud (VPC)

Availability Zone

Availability Zone

Availability Zone

ECS Cluster




Auto Scaling group



t3.medium



t3.medium



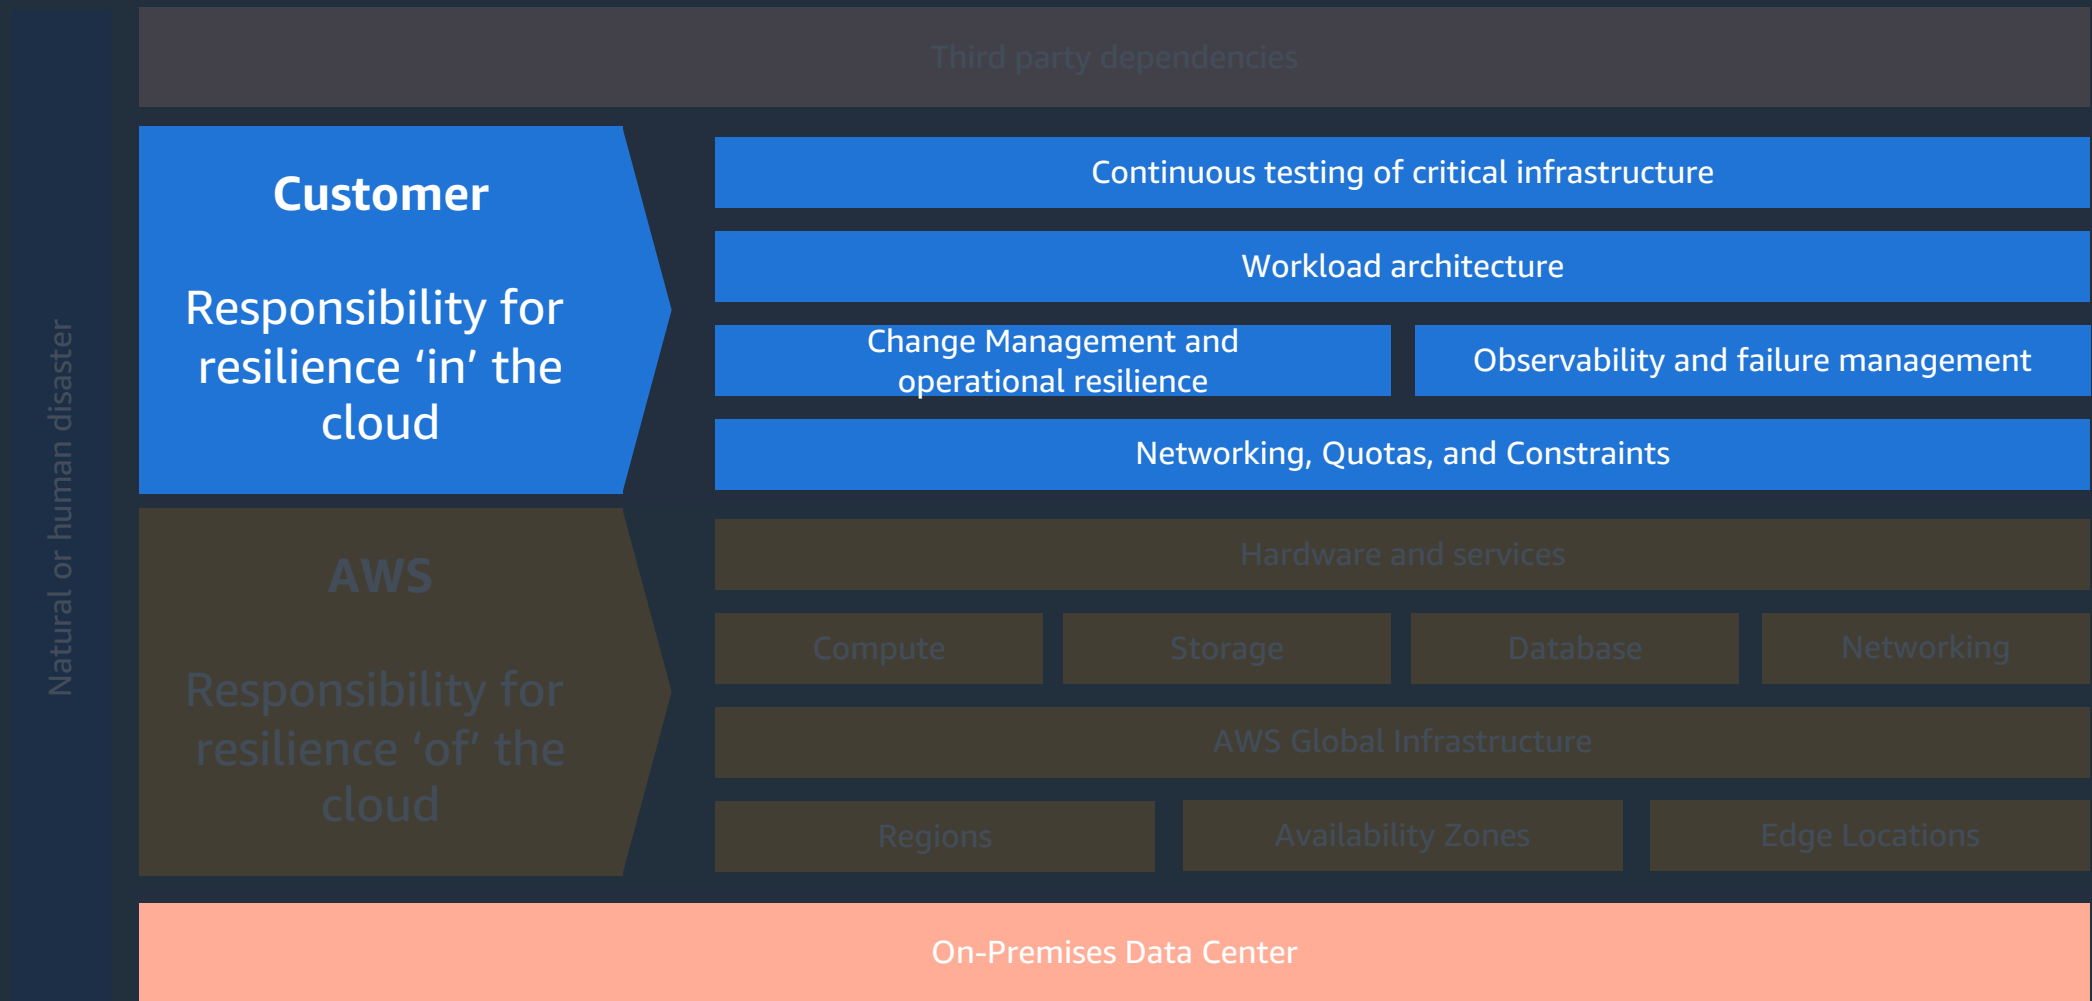
t3.medium



Let's formulate an experiment

- **Steady State** Application is reachable. Response time is <500ms
- **Hypothesis** What if a ECS node gets terminated?
- **Experiment** Turn of a EC2 instance running the ECS cluster

Resilience on AWS – A shared responsibility





Thank you!

Marcel Neidinger

mneiding@amazon.com

A few things you can do

DDoS yourself*

```
$ wrk -t12 -c400 -d30s http://127.0.0.1/api/health
```

Run CPU Stress with stress-ng

```
$ stress-ng --cpu 0 --cpu-method matrixprod -t 60s
```

Add latency to your network

```
$ tc qdisc add dev eth0 root netem delay 300ms
```

